SPERT® Forecaster

Quick Reference Guide

Agile Release Forecasting with Monte Carlo Simulation · forecaster.spertsuite.com

WHAT IS SPERT® FORECASTER?

Sprint velocity	The amount of work a team completes per sprint (in any unit: story points, hours, items). SPERT® Forecaster derives velocity mean and standard deviation from your sprint history, or accepts a subjective estimate.
Remaining backlog	The total work yet to be done at the time of the forecast. Enter the current backlog size in the same unit as your velocity.
Distribution types	Six distributions are available: T-Normal (Truncated Normal), Lognormal, Gamma, Triangular, Bootstrap (history mode, 5+ sprints), and Uniform (subjective mode). Only Lognormal is enabled by default — add others in Settings → Statistical methods to show. Each enabled distribution produces its own set of percentile results.
Percentile results	P50 = 50% chance of finishing by that sprint/date. P80 = 80% confidence. Higher percentiles give later dates with more buffer. Use P80–P90 for stakeholder commitments.
Forecast modes	History — derives the velocity mean and variability from your recorded sprints. Subjective — uses a point estimate plus a coefficient of variation (CV) when history is thin. **The app chooses for you:** with 5 or more included sprints it starts in History, otherwise Subjective. You can switch to History manually from 2 included sprints onward, provided those sprints do not all share an identical velocity — zero variability would imply a certainty the data cannot support, so the app disables the History button in that case.

■ SPERT® Forecaster can take a product across from SPERT® Story Map — either as an exported file, or handed straight over by Story Map's **Send to SPERT Forecaster** button — converting releases into milestones with pre-populated backlog sizes and velocity data. Re-importing the same product can merge into your forecasting setup rather than overwrite it; see IMPORTING FROM SPERT® STORY MAP.

STEP-BY-STEP WORKFLOW

1. **Projects** — Click **New Project** and enter a name and a unit of measure (e.g., story points, hours); start and finish dates are optional. Or click **Load Sample** for a ready-made example with sprint history, milestones, and a productivity adjustment. Load Sample is available at any time from the toolbar — not only when your workspace is empty — so you can pull a fresh reference copy alongside your own work. Each project is a tile: click anywhere on it to open its Sprint History. **Delete** sits beside the name; point at a tile, or tab into it, to show its other buttons — **Export** (that project as JSON), **Edit**, **Clone** (copies all data to a new project), **Archive** and **Tile color**.

2. **Sprint History** — First set the **Sprint Cadence** (1–4 weeks) and **First Sprint Start Date**: the forecast needs both. Then record completed sprints. For each sprint, enter the *Done this sprint* value (velocity). Optionally enter the backlog remaining at sprint end for burn-up chart accuracy. A compact read-only summary of your 3 most recent sprints appears above the Add Sprint form for quick reference. Toggle sprints in/out of the forecast with the Include checkbox to exclude anomalous sprints.

3. **Forecast tab** — Select History or Subjective mode. Enter the remaining backlog. Velocity mean and variability are auto-populated from history. Once sprint cadence and start date are configured, the simulation runs automatically on first valid input — no manual click required while Auto-recalculate is on (the default). If a prerequisite is missing, a helper message under the Run Forecast button explains what to set on the Sprint History tab.

4 **Adjust inputs** — Optionally: adjust velocity volatility using the Adjust link (history mode), or click *Reset overrides* to clear all manual velocity and variability overrides at once. Model scope growth from the calculated or custom rate, or switch to Subjective mode to enter a point estimate and predictability rating.

5 **Review charts** — The "Your forecast" hero callout shows the true cumulative probability at the displayed date (capped at 99%). Use the scope picker to switch between Entire Project and individual milestones. The Forecast Results table is collapsed by default — click its header to expand. The **Deadline Probability** panel (below the results table) answers the inverse question: given a target date, what is the probability of finishing by then? Scroll to explore: Burn-Up, CDF, and Histogram charts. Open **Explore a custom percentile** for any P1–P99 result. Leaving the Forecast tab and coming back no longer discards your results — you get the same dates back rather than a fresh simulation. Only the latest run is kept: forecast another project, and this one needs a fresh run when you return.

6 **Share results** — Use the copy-image button on any chart to capture it to the clipboard, or click the report icon to generate a printable HTML report including the results table and selected charts. Use milestones to show per-release forecast dates.

■ Auto-recalculate (Settings) defaults to ON — the simulation runs automatically whenever inputs are valid. For large trial counts or slow devices, disable auto-recalculate and click Run Forecast manually.

FORECAST INPUTS — FIELD GUIDE

Backlog	Remaining work to complete (required). Shown in the project's unit of measure. The helper text (*Last sprint: N*) shows the Backlog at End of the most recent included sprint that has one, which also fills the field. If the stored value drifts from the most recent *included* sprint's ending backlog (e.g., after toggling a sprint's Include status), a *Reset to N* link appears next to the field for one-click recovery. Your entries on this tab are kept for each project while the page is open, but not across a reload.
Velocity (mean)	Average sprint throughput. Auto-populated from history; overridable. In Subjective mode, derived from your point estimate — the field is read-only.
Variability	Velocity spread (standard deviation). A help tooltip on the label reads: "Standard deviation — how much velocity varies sprint to sprint." Auto-calculated from history. Use the *Adjust* link to apply a volatility multiplier (0.75×, 1.0×, 1.25×, 1.5×) instead of entering a raw value. In Subjective mode, derived from the selected CV. When any velocity or variability override is active, a *Reset overrides* link clears all three fields and closes the volatility adjuster in one click.
Start Date	The forecast start date — automatically set to the next sprint's start date based on your sprint cadence and history. For a project that has a first sprint start date set but no completed sprints yet, the forecast anchors on that *first sprint's* start date rather than today, so every derived date (results, CDF, burn-up, deadline probability) is computed from the schedule you entered. Read-only; update via the Sprint History tab.
Cadence	Sprint length in weeks (1–4). Set on the Sprint History tab, with the first sprint start date; displayed here for reference. Affects how sprint counts translate to calendar dates in all results. Sprint cadence and first sprint start date must be configured before the Run Forecast button is enabled, and every stored sprint date must be one the app can read: if one is not, the Forecast tab says why and offers **Go to Sprint History**.
Scope growth	Optional; appears once at least 2 sprints have a Backlog at End. Model backlog expansion sprint-over-sprint (if backlog tends to grow each sprint). Use the calculated rate (derived from Sprint History backlog entries) or enter a custom value. Adds the growth rate to remaining backlog each simulated sprint. The *Model scope growth* toggle is remembered per project — switching projects no longer carries one project's setting into another.

Subjective mode	Enter an estimated velocity and select a predictability rating: **Very steady** (CV 0.15) → **Fairly consistent** (0.25) → **Somewhat variable** (0.35) → **Often disrupted** (0.45) → **Quite volatile** (0.55) → **Wildly uncertain** (CV 0.65). Std Dev = velocity × CV. This is the mode the app starts in until 5 sprints are included, and it stays the right choice whenever your history is too short or too uniform to describe future variability. If your estimate is more than 2× or less than 0.5× the historical average, an amber warning appears — verify the estimate is intentional.

FORECAST RESULTS & CHARTS

Forecast Summary	A blue "Your forecast" hero callout shows the *true cumulative probability* at the displayed sprint-end date — typically higher than the selected percentile because sprint quantization means each sprint bucket covers a range of simulated outcomes. The probability shown in the callout is capped at 99%; the CDF chart is not capped and correctly reaches 100%. The scope picker lets you switch between Entire Project and individual milestones; both the hero callout and the per-milestone breakdown update immediately. The summary sentence reads *"there is at least an X% chance"* — the "at least" reflects the sprint-quantization gap. A colored-dot list below shows forecast dates for each milestone at the selected percentile. Completed milestones appear in italic. Leave the Forecast tab and return and the same numbers come back, rather than a re-run producing slightly different dates; the summary's scope, distribution and percentile selections are remembered for each project.
Percentile table	Rows = selectable percentiles (P5–P95 toggleable; default is P10, P50, P90); columns = enabled distribution types. Each cell shows the sprint the work finishes in (counting sprints already done) and that sprint's finish date. The Forecast Results section is **collapsed by default** — click its header to expand. The summary's percentile dropdown spans P10–P95; for any other value, open **Explore a custom percentile**, which has two sliders, P1–P99 (P85 and P50 unless you change the defaults in Settings). In History mode with fewer than 4 included sprints, an amber warning notes that forecast spread may be understated. Your percentile pill selection and both custom-percentile slots are remembered per project; they no longer re-seed from Settings each time you return to the tab.
Deadline Probability	A collapsible panel (collapsed by default) that inverts the forecast question: given a target date, what is the probability of finishing by then? Set a Target Date, choose a Scope (Entire Project or any not-yet-completed milestone), and select a Distribution. The panel shows a per-distribution probability table and a narrative sentence naming the scope, the date, and the sprint at which the probability is evaluated. Probability is measured at the end of the last forecast sprint whose finish date falls on or before your target date (conservative: "by this date" = "by this sprint boundary"). A sprint-boundary footnote always names the exact sprint and date used so you can audit the result. All probabilities cap at 99%. The target date is kept for each project until you reload the page.
Milestone results	If milestones are defined, the Forecast Summary lists each one's forecast date at the selected percentile, and Forecast Results shows a table for each upcoming milestone. The Milestone dropdown on the CDF chart, the Histogram and Explore a custom percentile focuses on one release; it offers milestones not yet completed whose Chart box is ticked. **Entire Project** always means the whole remaining backlog, never a dropdown's choice: when your milestones add up to less than the backlog, the dropdowns end with *Entire Project (Total)* and Forecast Results adds an Entire Project table.
Burn-Up chart	Plots completed work and total backlog scope (solid lines; *Show scope* hides the scope line), three dashed forecast lines — by default Optimistic (P10), Expected (P50) and Conservative (P90) on Lognormal, each with a label, percentile and color you can change — and a colored reference line for each charted milestone. Chart legend is ordered: Scope, Done, then forecast lines ascending by percentile. Hover for sprint details.

CDF chart	Cumulative distribution function — probability of finishing within each number of sprints. Shows all enabled distributions as overlapping curves. Hover to read exact probabilities. A dashed line marks your custom percentile. The Milestone dropdown chooses which milestone the curves show; completed milestones are left out of it.
Histogram	Frequency distribution of simulated sprint-to-finish counts across all trials. Visualizes the spread of possible outcomes — a narrow peak = low uncertainty; a wide spread = high volatility. Completed milestones are excluded from the milestone selector dropdown.
Copy & Report	Each chart has a copy-image button (two overlapping squares) to capture it to the clipboard; Firefox does not support it. The report icon in the Forecast Results header (expand the section to see it) opens a popover to generate a printable HTML report including the results table and any expanded charts. All chart captures render correctly in both light and dark mode.

■ The Bootstrap distribution (history mode, 5+ sprints; tick it in Settings → Statistical methods to show) resamples actual sprint velocities directly — no distributional assumption. It often gives the most realistic spread when you have a reliable history. In History mode with fewer than 5 included sprints, a footnote below the results table shows how many more sprints are needed to unlock Bootstrap.

MILESTONES & PRODUCTIVITY ADJUSTMENTS

Milestones	Named delivery targets (e.g., MVP, Beta, GA). Each milestone's **Remaining Work** is incremental — the work still to do in that milestone only. Update it as work is done; 0 marks the milestone completed. Milestones are cumulative: the forecast chains them sequentially. You can add up to 10 milestones by hand; Story Map updates can bring in more, and this app's own files hold up to 100 per project. **Drag to reorder:** a blue insertion line marks the gap the milestone will land in and follows your pointer within a row — hold a row's top half to target the gap above it, its bottom half for the gap below — so every gap is reachable from either direction, including the one beneath the last milestone. To fix a name, click it directly on the Milestones panel — the cell becomes a text input in place. Use the Pencil button when you also need to change Remaining Work or Color.
Milestone colors	Each milestone gets an auto-assigned color (emerald, blue, amber, purple, red — cycled). The color marks the milestone's reference line on the Burn-Up chart and its dot in the forecast summary. Untick its **Chart** box to take it off the Burn-Up chart and out of the milestone dropdowns and results tables.
Productivity adjustments	**Productivity Adjustments (Holidays, Breaks, Events)** — Define date ranges where the team operates at reduced capacity. Enter a start date, end date, and productivity factor: **0%** = no work done, **50%** = half velocity, **100%** = full velocity (no effect). Enabled adjustments reduce effective velocity in the simulation for those sprints. Note: because forecasts report sprint finish dates (not intra-sprint completion dates), a small adjustment may not shift the projected end date if work still falls within the same sprint. An adjustment that ends before the forecast's first sprint has no effect.
Enable / disable	Each productivity adjustment has a toggle. Disabled adjustments are saved but have no effect on the forecast — useful for modeling scenarios without deleting the adjustment.

■ Milestone backlog sizes are incremental, not cumulative. If your total backlog is 500 points split across MVP (200) and GA (300), enter 200 for MVP and 300 for GA — not 200 and 500. SPERT® Forecaster accumulates them automatically. If they add up to less than the remaining backlog, Entire Project still forecasts the whole backlog; if they add up to more, the summary says so, and any milestone past the backlog is dated at project completion.

SPRINT HISTORY TAB	
Sprint config	Set the sprint cadence (1–4 weeks) and first sprint start date here before adding a sprint; both lock once a sprint is recorded. Project start and finish dates are set in the project dialog on the Projects tab. Sprint start dates are calculated automatically from the first sprint start date and cadence.
Adding sprints	Click Add Sprint. A compact read-only summary of your 3 most recent sprints appears above the form — showing sprint number, date range, velocity, and backlog at end — so you can reference prior values without closing the form. Enter the *Done this sprint* value (velocity) and optional *Backlog at End* (in the project's unit of measure). Sprint finish dates are calculated and shown; optionally override a finish date to extend a sprint (e.g., for a holiday week or spring break). Only the most recent sprint can be deleted.
Custom finish dates	Override the computed finish date for any sprint when the team worked longer than the standard cadence. Overridden dates cascade forward — all subsequent sprint start dates shift accordingly. A visual indicator marks sprints with custom dates. Use Reset to restore the calculated date. A stored date the app cannot read is flagged with a warning in its row and a notice saying how to fix it; until then the forecast will not run.
Include toggle	Exclude anomalous sprints (team at 20% due to vacation, major scope change, etc.) from velocity statistics without deleting the record. Excluded sprints still appear in the burn-up chart. If the Backlog field on the Forecast tab drifts after a toggle, a *Reset to N* link appears next to the field for one-click correction.
Velocity stats	Included Sprints, Velocity (the mean) and Variability (the standard deviation) are shown above the sprint list. These feed directly into the Forecast tab.
Velocity chart	**Velocity Trend** (collapsed; appears once there are 2 sprints): a line chart of per-sprint velocity with a dashed mean line, a trend line, and a label saying whether throughput is Improving, Declining or Stable. Visually identify trends and outliers over time.
Scope analysis	If Backlog at End is recorded for 2 or more sprints, the **Scope Change Analysis** panel (collapsed) shows how the backlog changed per sprint. The same figures feed the Forecast tab's calculated scope growth, which adds back each sprint's completed work to measure new scope.

DATA, SETTINGS & STORAGE	
Local storage	Default. Data stays in your browser only — never sent to any server, unless you turn on Connect AI's Read Mode (see CONNECT AI). Safe for corporate data. Clearing browser site data permanently deletes all projects. Use Export to back up.
Cloud storage	Optional. Sign in with Google or Microsoft. Data syncs via Firebase across devices. Enables project sharing with team members (owner / editor / viewer roles). **Microsoft sign-in requires a work or school account** — personal Microsoft accounts (outlook.com, hotmail.com, live.com) are refused before a password is entered. Microsoft enforces this, not the app, so it applies however sign-in is reached. Google is unchanged and still accepts personal accounts. In cloud mode the Projects, Sprint History and Forecast tabs show *Loading your projects from the cloud…* until your projects arrive; if they cannot load, **Try again** and **Export a backup** are offered. A project shared with you as a viewer is marked **View only**, and one this browser kept that is no longer in your cloud account **Not in cloud**: you can forecast and export it, but not change it. Only a project's owner can delete it. Each device keeps its own project order.

Export / Import	**Export** in a tile's row of buttons exports that single project as JSON. The **Export All** button in the toolbar exports every project; Settings → Export Projects lets you select a subset. Every JSON export checks that its file will import again; if it will not, a warning — *Your file was saved, but it will not restore.* — names each project and its problem. When a file holds a project you already have, **Import** opens an inline conflict-resolution preview: each conflicting project shows three choices — *Keep existing* (default when the identifiers match), *Add as a copy* (default when only the names match), or *Replace existing* — and files from SPERT® Story Map add a fourth, *Update* (next section). A project you can only view offers Keep or Copy alone. A copy is named *Name - Copy (N)*, matching the Clone button convention. Non-conflicting projects are listed in a green block as new. A full backup from Export All opens on **Replace all data** when you already have projects, which replaces every one of them; choose **Merge into workspace** to keep them and decide project by project. A result banner confirms counts after the import completes, and both the preview and that banner now survive leaving the Projects tab and coming back. Imports are capped at 10 MB. In cloud mode the Import button stays disabled until your cloud projects have finished loading, so conflict detection always runs against your real workspace.
Archive & tile color	**Archive** takes a finished project off the dashboard and out of the Sprint History and Forecast project lists. Tick **Show archived (N)** under the heading to see archived tiles; click one to unarchive it and open its Sprint History, or use its Unarchive button. With every project archived, Sprint History and Forecast say so and the Connect AI button is hidden. **Tile color** gives a tile a light wash: Blue, Teal, Slate, Purple, Pink or No color. Both are your own settings: they never change a shared project for anyone else, and exported files carry neither.
Import from Story Map	(Projects tab → Import, or **Send to SPERT Forecaster** from inside SPERT® Story Map.) Releases become milestones; sprint velocity and backlog data are mapped automatically, and both routes use the same conflict-resolution preview as native Forecaster exports. Re-importing a product you already hold can *Update* it in place, keeping your forecasting setup. The next section covers both routes and the merge rules in full.
Export CSV	The CSV button in the Forecast Results header downloads the simulation behind your results: its parameters, productivity adjustments and milestones, the percentile results, the frequency distribution, and every trial — for reports or your own analysis.
Settings	Configure: storage (Local or Cloud; sign in under Account), trial count (1K–50K; default 10,000), auto-recalculate toggle (defaults to ON — and it now skips the re-run entirely when nothing affecting the numbers has changed), **Statistical methods to show** (enable/disable individual distributions; default is Lognormal only; at least one must remain enabled), the startup warning when using local storage, default chart font size, default custom percentiles (two slots), default results percentiles, theme (light/dark/system), and Export Attribution (name and ID embedded in exported JSON for traceability).

■ Set a lower trial count (1,000–5,000) for fast exploratory forecasting, then increase to 10,000–50,000 before a stakeholder presentation for smoother, more stable percentile results.

IMPORTING FROM SPERT® STORY MAP

Two routes	**Export a file** from SPERT® Story Map and Import it here — the only route that works between two different machines, or that leaves you a copy to keep. Or use **Send to SPERT Forecaster** inside Story Map, which opens this app in a new tab and hands the project over directly, with no file on disk. The handover works only between two tabs in the same browser and is a one-time transfer: nothing stays attached, nothing polls, and closing either tab ends it. What arrives is byte-for-byte what the file would have carried, and both routes run the same validator, the same 10 MB ceiling and the same conflict preview. In cloud mode a transfer arriving while your cloud projects are still loading is not refused: its preview waits until they have loaded, and nothing is imported until you press **Apply import**.

Conflict preview	When an incoming project matches one you already have, an inline preview opens and **nothing is written until you press Apply import**. Each conflicting project offers *Keep existing, ignore imported*; *Add as a copy*; *Replace existing with imported*; and — for Story Map files only — *Update with new progress, keep my forecast setup*. Non-conflicting projects are listed separately as new. A project you can only view offers the first two alone.
What Update keeps	Update merges rather than substituting. Sprint progress and any new sprints come from Story Map, while everything you configured here stays: project start and finish dates, productivity adjustments, custom sprint finish dates, your sprint-exclusion choices, burn-up chart configurations, unit of measure, and each milestone's color and chart setting. From a Story Map file older than v0.53.8 it keeps your milestone figures too. Update is the pre-selected choice wherever it is offered.
What Update overwrites	Story Map is treated as the authority on progress and sequence, so anything it also tracks is overwritten on *every* refresh: a completed-work figure you corrected by hand, a project or milestone you renamed here, any milestone reordering you did here, and — from Story Map v0.53.8 or later — each matching milestone's remaining work, including a figure you changed here. If you rely on one of those, Update will undo it each time — Replace and Add as a copy remain available.
When Update is withheld	Update needs evidence that the incoming project *is* the one you already have. A matching identifier is that evidence, and so is at least one shared sprint — which is what keeps Update reachable for a project whose identifier changed during a move to cloud storage. Two projects that merely share a name, with no sprint in common, still get Skip / Copy / Replace only. Update is also withheld when your copy holds a sprint the incoming file does not, because Story Map renumbers sprints on every export and merging would silently move the forecast by a full cadence period. The screen states the reason and names both ways forward: delete that sprint on the Sprint History tab, or choose Replace.
Milestone figures and ordering	Story Map v0.53.8 and later send each release's **work remaining**, which is what a milestone's figure means here, and say so in the file. Older Story Map files send each release's **total scope** in the same field: from those, a milestone arriving for the first time carries a total where this app expects a remainder, and reads high until you correct it. Milestones you created here can never match a Story Map release, so they are always kept — and placed *after* the imported ones. Because milestone figures add up in order, a kept milestone's target then includes every imported milestone ahead of it, and its forecast date follows their figures.
After the import	A summary banner reports what was actually written rather than what was previewed: sprints refreshed and added; each milestone figure Update replaced, old → new, and any milestone that completed or reopened; any milestones added from Story Map, named individually (from an older file, with the total-scope warning above); any milestones kept that Story Map did not send; and the ordering effect. An archived project stays archived, and the banner says so. Your forecast *run* is cleared — the sprint set moved, so the old numbers no longer describe it — while your forecast *view* (target date, selected milestone, scope-growth settings) is kept, which is the difference between Update and Replace. With auto-recalculate on, a fresh run starts by itself.

■ The intended rhythm is a round trip: story-map the release in SPERT® Story Map and complete sprints there, then send or export to the Forecaster to see when it lands. Update exists for the second crossing and every one after it — it is the only choice that takes the new progress without destroying the forecasting setup you built here. Skip discards the progress you just exported; Replace takes the progress and clears your dates, holiday and vacation periods, sprint exclusions, custom finish dates and burn-up configuration.

■ **If your sprint history vanished after you switched to cloud storage, this is why.** A project whose internal identifier had to change during the upload left its sprints behind, and the upload still reported success — the loss only showed on the Sprint History or Forecast tab, by which point the local copy was gone. Migration runs one way, so there is nothing to roll back: re-import the project from its original file or from SPERT Story Map to restore it. Fixed in v0.42.1 — sprints now move with the project, and a migration that assigns a new identifier says so.

CONNECT AI	
Connect AI button	Pair an AI assistant — Claude, ChatGPT, Copilot, Gemini and others — with the project you have open, so it can read your forecast and explain it. The button sits in the app header, immediately left of the theme toggle, and is visible from every tab. Once paired it reads **AI** with a status dot: **grey** means an assistant is paired but idle, **pulsing blue** means one is actually connected and reading. The button is hidden until you have at least one project that is not archived, and requires cloud configuration — it does not appear in a local-only build.
Read-only, always	The connection cannot change anything in SPERT® Forecaster. There is no tool an assistant could call to edit your project — the server does not define one — the pairing is created with writes permanently refused, and the server re-checks that before any write could occur. If your assistant concludes something should change, it will tell you what and where, and you make the change yourself.
Consent & Read Mode	The first connection opens a consent dialog. **Read Mode** is the only toggle: with it on, a copy of the *currently open* project is uploaded to SPERT's cloud so the assistant can read it. Only that one project is ever uploaded — never your whole workspace — and switching projects replaces the copy rather than accumulating them. An archived project is never the open one: archive it and the assistant follows the first project you have not archived; archive every project and the uploaded copy is deleted and the button hides, though the pairing is not ended. Turning Read Mode off deletes the uploaded copy immediately and keeps the pairing alive; with it off, the assistant can still explain how forecasting works but cannot see your data. Once you have consented, later connections from the same browser resume silently and go straight to the pairing code, until you Disconnect or sign out; after that, the consent dialog opens again.
Pairing code & prompt	The panel shows a single-use pairing code that expires in 15 minutes (it refreshes every 10 minutes while the panel is open and waiting). Two buttons copy what you need: **Copy code** for the code alone, and **Copy prompt for your AI** for a prepared prompt that carries the code and tells your assistant how to read the data correctly — use the prompt unless you have a reason not to. Your assistant needs the SPERT MCP server (mcp.spertsuite.com) added as a connector first. A pairing stays valid for 7 days from your last activity.
What the assistant reads	Your project configuration, sprint history and velocity statistics, the inputs the forecast actually ran on, percentile results per distribution and per milestone, your productivity adjustments, and your deadline probability. It does *not* receive the raw per-trial data, the charts as drawn, any other project, the cumulative-probability curve, or the large percentage in the "Your forecast" headline — that figure is a different calculation and reads higher than the percentile below it. The data carries its own list of what it omits, so the assistant can say what it cannot see rather than guessing.
Freshness status	Every read is labelled, so an assistant need not guess whether your numbers are current: **fresh** (the results match your current inputs), **stale** together with the specific reason — a milestone was resized, the included sprint set changed, a productivity adjustment moved — **recomputing** (a run is in progress), or **absent** (no forecast has been run in this browser session, or none can run because an input is missing or a sprint date is invalid). A project too large to send is marked **unavailable** and carries its settings only. Ask your assistant to check the status before it quotes a date.
Visible vs. computed	The app computes up to six distributions but you may be looking at one. The data says which distributions are actually visible to you as distinct from which were computed, and separates the ones excluded because they are invalid for your current mode from the ones you chose to hide in Settings. Each scope is likewise flagged with whether it is the one on your screen, so an assistant will not quote a milestone you are not reading.

Disconnect	**Disconnect** in the panel ends the pairing and deletes both the session and the uploaded copy. Signing out does the same, and so does signing in as a different account. If signing out cannot reach the server, a notice says your assistant may still be connected and offers **Try again**; the app also retries by itself, and the session ends on its own within 7 days. Turning Read Mode off deletes the uploaded copy but keeps the pairing, and switching from cloud to local storage does likewise. The panel also closes on **Done**, on Escape, or on a click outside it.

■ Tell your assistant this is a **Monte Carlo simulation over sprint velocity** — it is **not** three-point PERT, and there is no (O + 4M + P) / 6 anywhere in it. Assistants get this wrong reliably and confidently. The copied prompt says so, and two explainer tools (available to any assistant connected to the SPERT MCP server, with or without a pairing) cover the forecasting method and this app's vocabulary. Ask your assistant to consult them before it explains the method.

■ Two things worth knowing. **Read Mode uploads regardless of your storage setting** — if you work in Local storage, enabling Read Mode still sends a copy of the open project to SPERT's cloud for as long as the pairing lasts. And a pairing left open for days with no edits can lose its uploaded copy to expiry while still looking connected; if your assistant reports no data on a healthy-looking pairing, make any edit, or toggle Read Mode off and back on.

KEYBOARD SHORTCUTS

1	Go to Projects tab
2	Go to Sprint History tab
3	Go to Forecast tab
4	Go to Settings tab
5	Go to About tab
?	Open the Keyboard Shortcuts help (Esc, its × button or a click outside closes it)

■ Shortcuts do nothing while you are typing in a field or have a dropdown, checkbox or date field selected, and while any dialog is open — the shortcuts help included. Press ? on any tab to see the shortcut list inside the app.

DISTRIBUTION TYPES AT A GLANCE

T-Normal	Truncated Normal — symmetric bell curve restricted to non-negative velocities; classical PERT-style shape. A reasonable alternative when sprint variability is low and approximately symmetric — at high variability the lower-bound truncation biases the mean upward and makes forecasts artificially optimistic. Available in both History and Subjective modes.
Lognormal	Lognormal — right-skewed curve always above zero, with a long upper tail; the recommended default because it matches the empirical right-skew of sprint velocity and remains well-calibrated even when sprint-to-sprint variation is large. **Enabled by default;** other distributions require opt-in via Settings → Statistical methods to show. Available in both History and Subjective modes.
Gamma	Gamma — right-skewed like Lognormal but with a thinner upper tail. Useful when faster sprints happen but extreme breakouts are unlikely. Available in both History and Subjective modes.
Bootstrap	Bootstrap — resamples directly from your actual sprint history; the most data-driven option, assuming only that future sprints will look like past ones. Requires 5+ included sprints (History mode only). Most reliable with 10+ representative sprints.
Triangular	Triangular — a simple peaked shape with hard limits at ±3 standard deviations from the mean, never below zero (at high variability the limits are zero and twice the mean). Useful when you want a transparent, bounded forecast without long tails. Available in both History and Subjective modes.

Uniform	Uniform — every velocity in the range equally likely; the most conservative shape. Useful when you have little basis to prefer any one velocity value. Subjective mode only.

■ When multiple distributions are enabled and agree on a percentile date, that consensus signals a robust forecast. Significant divergence indicates sparse or volatile velocity data — consider enabling additional distributions in Settings → Statistical methods to show, and gathering more sprint history before committing to a deadline.

FORECAST GUIDANCE FOR NEW USERS

Choosing a percentile	Use **P50** for internal team targets — equal odds of finishing early or late. Use **P80** for stakeholder commitments — 4 in 5 chance of on-time or better delivery; this is the most common choice for sprint-based Agile releases. Use **P90–P95** for contracts, hard external deadlines, or releases with downstream dependencies. Avoid anchoring on P99 — it adds excessive buffer and rarely improves stakeholder trust.
Sprint count thresholds	**0–1 sprints:** Subjective mode only — History mode requires 2+ included sprints. **2–4 sprints:** History mode available; T-Normal, Lognormal, Gamma, and Triangular can be shown (tick them in Settings); with 1–3 included sprints an amber warning notes that forecast spread may be understated. **5+ sprints:** the app starts in History mode, and the Bootstrap distribution unlocks (Uniform is Subjective mode only). **10+ sprints:** Distributions typically converge — most reliable results. The more representative sprint history you have, the narrower and more trustworthy the forecast spread.
Forecast date precision	Projections show **sprint finish dates**, not specific days within a sprint. A productivity adjustment that reduces velocity by a small amount may not shift the projected end date if the adjusted velocity still completes the remaining work within the same sprint — this is correct behavior, not a bug. To model a full sprint gap (e.g., a team fully offline for a sprint), set the productivity factor to 0% for that date range, or, for a sprint already recorded, extend it with a custom finish date.
Backlog field units	The *Backlog* field on the Forecast tab and the *Backlog at End* field on the Sprint History tab must be entered in the same unit of measure configured for the project — the same unit used to record sprint velocity. Example: if the project is set up in story points and the team has 41 backlog items worth roughly 1,100 story points, enter **1,100** — not 41. Entering the raw item count when the project unit is story points will dramatically underestimate scope and produce an optimistic but incorrect forecast. If the project unit is *items*, then entering the item count is correct.
Reading chart signals	**Histogram:** a narrow peak = consistent team velocity and a tight forecast; a wide spread = high variability; a long right tail = some trials finishing much later, from runs of slow sprints. **CDF:** a steep S-curve = low variance and a tight forecast; a gradual slope = wide uncertainty; a nearly vertical step with nearly identical percentile dates suggests all sprints had the same velocity or you have too few data points — add more sprint history before committing to a deadline.

■ If all distribution columns in the percentile table show the same completion date, you likely have identical sprint velocities (std dev = 0), or a remaining backlog small enough to finish in the same sprint in almost every trial. Switch to Subjective mode and select a CV that reflects your team's expected variability, or add more sprint history before making stakeholder commitments.

SPERT® Forecaster · © 2026 William W. Davis, MSPM, PMP · Licensed under GNU GPL v3 forecaster.spertsuite.com